

COSC 4740
Final Project

Due date: December 8, 2015

I hear and I forget, I see and I remember, I do and I understand. - Chinese proverb

Goal: To design and implement a subset of the Unix file system.
Grade: 30% of your total grade is allocated for this project.

The filesystem consists of five classes:

1. *Disk*: Simulates the behavior of a disk that is used to store files and other filesystem information.
2. *DiskManager*: Partitions and manages access to a disk.
3. *PartitionManager*: Manages partitions in a disk.
4. *Filesystem*: Implements various filesystem operations.
5. *Client*: Uses a filesystem.

Disk

Simulates the behavior of a disk using a Unix file. It exports the following functions:

```
int initDisk()
int readDiskBlock(int blknum, char *blkdata)
int writeDiskBlock(int blknum, char *blkdata)
int getBlockSize()
int getBlockCount()
```

The first function creates a Unix file to simulate a disk. If such a file already exists, it returns 0, otherwise it creates the file of specified size and returns 1. The second function reads the disk block numbered *blknum* into the buffer pointed to by *blkdata*. The third function copies the buffer pointed to by *blkdata* to the disk block numbered *blknum*. Both of these functions return 0 if successful and a negative number otherwise. Every Disk object is managed by a DiskManager object.

DiskManager

Partitions and manages access to a disk. There is one DiskManager object associated with each Disk object. A DiskManager object partitions a disk into multiple partitions, and provides some of the basic functions needed by a filesystem to store its information. These functions are:

```
int readDiskBlock(char partitionname, int blknum, char *blkdata)
int writeDiskBlock(char partitionname, int blknum, char *blkdata)
int getBlockSize()
int getPartitionSize(char partitionname)
```

readDiskBlock reads the partition block numbered *blknum* into the buffer pointed to by *blkdata*. *writeDiskBlock* copies the buffer pointed to by *blkdata* to the partition block numbered *blknum*. *getBlockSize* return the size of a disk block

in bytes. `getPartitionSize` returns the total number of blocks available in this partition.

PartitionManager

A `PartitionManager` object manages partitions in a disk. A `PartitionManager` object is created by a `FileSystem` object (described next) and has a partition name associated with it. Each `PartitionManager` object is associated with a `DiskManager` object. Each `PartitionManager` has a bitmap, that keeps track of free and allocated blocks. The bitmap is written out block 0 of the partition. It exports the following operations:

int readDiskBlock(int *blknum*, char **blkdata*)

`readDiskBlock` reads the block numbered *blknum* in a disk partition into the buffer pointed to by *blkdata*.

int writeDiskBlock(int *blknum*, char **blkdata*)

This operation copies the buffer pointed to by *blkdata* to the block numbered *blknum* in a disk partition.

int getBlockSize()

This operation returns the size of a disk block in bytes.

int getFreeDiskBlock()

This operation allocates a free disk block in a partition. It returns -1 if there is no block available (ie the partition is full) or the block number allocated.

int returnDiskBlock(int *blknum*)

This operation deallocates a disk block in a partition. It will also write a blank (all c's) to the block that is returned. It returns 0 if successful or -1 for any other reason.

FileSystem

A `FileSystem` object implements a Unix file system. Each `FileSystem` object is associated with a `DiskManager` object, which it uses for all disk I/O. Multiple `FileSystem` objects may be associated with a single `DiskManager` object, which implies that multiple file systems can share a disk. This is equivalent to partitioning of a single disk drive into multiple directories in modern systems. Each partition is managed by a `PartitionManager` object, which is created by a `FileSystem` object.

Files and directories are implemented using a structure similar to the Unix i-nodes. There are three i-nodes: File i-node, directory i-node, and indirect inode. Every file in the file system has an file i-node associated with it. An file i-node will contain the following information: file name (1 byte), a 1 byte variable to indicate whether it is a file or a directory, file size (4 bytes), three disk block addresses (4 bytes each), and one single indirect block address (4 bytes). Use the remaining bytes in the block to store some other file attributes of your choice. An indirect block is 64 bytes (same as block size), with 16 spots for disk block addresses (4 bytes each). A directory i-node has the following structure, 64 bytes, with 10 slots containing: name (1 byte), block pointer (4 bytes), type file/directory (1 byte) and the last 4 bytes a pointer to the next block with the directory information will continue if needed. A `FileSystem` object exports the following operations:

int createFile(char **filename*, int *fname_len*)

This operation creates a new file whose name is pointed to by *filename* of size *fname_len* characters. File names and directory names start with ``/'` character and consist of a sequence of alternating ``/'` and alphabet (``A'` to ``Z'` and ``a'` to ``z'`) characters ending with an alphabet character. The `CreateFile` function returns -1 if the file already exists, -2 if there is not enough disk space, -3 if invalid filename, -4 if the file cannot be created for some other reason, and 0 if the file is created successfully.

```
int createDirectory(char *dirname, int dnameLen)
```

This operation creates a new directory whose name is pointed to by *dirname*. This function returns -1 if the directory already exists, -2 if there is not enough disk space, -3 if invalid directory name, -4 if the directory cannot be created for some other reason, and 0 if the directory is created successfully.

```
int lockFile(char *filename, int fname_len)
```

This operation locks a file. A file cannot be locked if (1) it doesn't exist, or (2) it is already locked, or (3) it is currently opened. It returns a number greater than 0 (lock id), if the file is successfully locked, -1 if the file is already locked, -2 if the file does not exist, -3 if it is currently opened, and -4 if the file cannot be locked for any other reason. A note, once a file is locked, it maybe only be opened with the lock id and the file cannot be deleted or renamed until the file is unlocked.

```
int unlockFile(char *filename, int fname_len, int lock_id)
```

This operation unlocks a file. The *lock_id* is the lock id returned by the `LockFile` function when the file was locked. The `UnlockFile` function returns 0 if successful, -1 if lock id is invalid, -2 for any other reason.

```
int deleteFile(char *filename, int fname_len)
```

This operation deletes the file whose name is pointed to by *filename*. A file that is currently in use (opened or locked by a client) cannot be deleted. It returns -1 if the file does not exist, -2 if the file is in use or locked, -3 if the file cannot be deleted for any other reason, and 0 if the file is deleted successfully.

```
int deleteDirectory(char *dirname, int dnameLen)
```

This operation deletes the directory whose name is pointed to by *dirname*. Only an empty directory can be deleted. This function returns -1 if the directory does not exist, -2 if the directory is not empty, -3 if the directory cannot be deleted for any other reason, and 0 if the directory is deleted successfully.

```
int openFile(char *filename, int fname_len, char mode, int lock_id)
```

This operation opens a file whose name is pointed to by *filename*. If *mode* = `'r'`, the file is opened for read only, If *mode* = `'w'`, the file is opened for write only, and if *mode* = `'m'`, the file is opened for read and write. An existing file cannot be opened if (1) the file is locked and *lock_id* doesn't match with *lock_id* returned by the `lockFile` function when the file was locked, or (2) the file is not locked and *lock_id* \neq -1. This operation returns -1 if the file does not exist, -2 if *mode* is invalid, -3 if the file cannot be opened

because of locking restrictions, -4 for any other reason, and a unique positive integer (file descriptor) if the file is opened successfully. If the file is opened successfully, an *rw* pointer (read-write pointer) is associated with this file descriptor. This *rw* pointer is used by some of the operations described later for determining the access point in a file. The initial value of an *rw* pointer is 0 (beginning of the file).

```
int closeFile(int filedesc)
```

This operation closes the file with file descriptor *filedesc*. It returns -1 if the file descriptor is invalid, -2 for any other reason, and 0 if successful.

```
int readFile(int filedesc, char *data, int length)  
int writeFile(int filedesc, char *data, int length)  
int appendFile(int filedesc, char *data, int length)
```

These operations perform read/write/append operations on a file whose file descriptor is *filedesc*. *length* is the number of bytes to be read from / written into / appended into the buffer pointed to by *data*. These operations return -1 if the file descriptor is invalid, -2 if *length* is negative, -3 if the operation is not permitted, and number of bytes read/written/appended if successful. The read and write operations operate from the byte pointed to by the *rw* pointer. The write operation overwrites the existing data in the file and may increase the size of the file. The append operation appends the data at the end of the file. The read operation may read less number of bytes than *length* if end of file is reached earlier. After the read/write/append is done, the *rw* pointer is updated to point to the byte following the last byte read/written/appended.

```
int seekFile(int filedesc, int offset, int flag)
```

This operation modifies the *rw* pointer of the file whose file descriptor is *filedesc*. The *rw* pointer is moved *offset* bytes forward if *flag* = 0. Otherwise, it is set to byte number *offset* in the file. This operation returns -1 if the file descriptor, *offset* or *flag* is invalid, -2 if an attempt to go outside the file bounds is made (end of file or beginning of file), and 0 if successful. A negative *offset* is valid.

```
int renameFile(char *fname1, int fname1_len, char *fname2, int fname2_len)
```

This operation changes the name of the file whose name is pointed to by *fname1* to the name pointed to by *fname2*. It returns -1 invalid filename, -2 if the file does not exist, -3 if there already exists a file whose name is the same as the name pointed to by *fname2*, -4 if file is opened or locked, -5 for any other reason, and 0 if successful. A note, you can rename a directory with this operation.

```
int getAttributes(char *filename, ...)  
int setAttributes(char *filename, ...)
```

These operations get/set the attributes of a file whose name is pointed to by *filename*. Work out the details of these operations based on the file attributes you choose. You must choose a minimum of **two** attributes for your filesystem.

Client

This class is used to create client objects that invoke file system operations. Each Client object is associated with a FileSystem object.

Implementation Guidelines

1. Implement your file system using C++.
2. Unix file system commands are used to implement the Disk. No other functions should use any Unix file system functions.
3. Use bitmaps to keep track of free disk blocks. The bitvector code is provided and is to be used to implement the bitmaps.
4. Some support files are provided to you to get started. These include `bitvector.h` and `bitvector.cc` that implement a `BitVector` class, `disk.h` and `disk.cc` that implement a `Disk` class, `diskmanager.h` and `diskmanager.cpp` that partially implement a `DiskManager` class, and `filesystem.h` and `filesystem.cc` class that partially implement a `FileSystem` class. In addition, an outline of a driver program to test the file system is provided in file `driver.cpp`. `Driver1.cpp` through `driver6.cpp` are used when turning in the project. All these files are available on the class web page.
5. Remember that your file system should survive the termination of clients that are using it.
6. Don't code until you understand what you are doing. Design, design, design first. While weeks of programming can save you hours of planning, I don't advise it.
7. Start working on this project now. First complete the implementation of `DiskManager` class.
8. Make sure you read everything, first! There is a lot of information in the `cpp` and header files and in other files provided.

Project Submission

1. You may do this project in teams of up to two students. However, you are responsible for submitting the final project irrespective of how your team partners work.
2. A note(s) page as the top page. This is detail how much of the project is completed and working. Including any information about the file attributes you choose. Say for example you only get to `driver4` completed and working, then you should state that.
3. If you can something about how this program compiles, then you need to document those changes in the `00_README` file and on the notes page.
4. Submit hardcopies of your source code and outputs of all 6 driver programs. Don't print the 6 driver files.
5. email me (nfraziel@uwyo.edu) the entire source (including drivers and any source code I have provided to you, `makefile`, `readmes`, etc) code of your program. Use *tar* and *gzip*. The subject line of the email should be 4740: Project

The last e-mail will be the only one accepted, so if you e-mail multiple copies (or multiple parts), I will only look at the last one and will only be accepted up to 5pm on December 8, 2015 as on time.